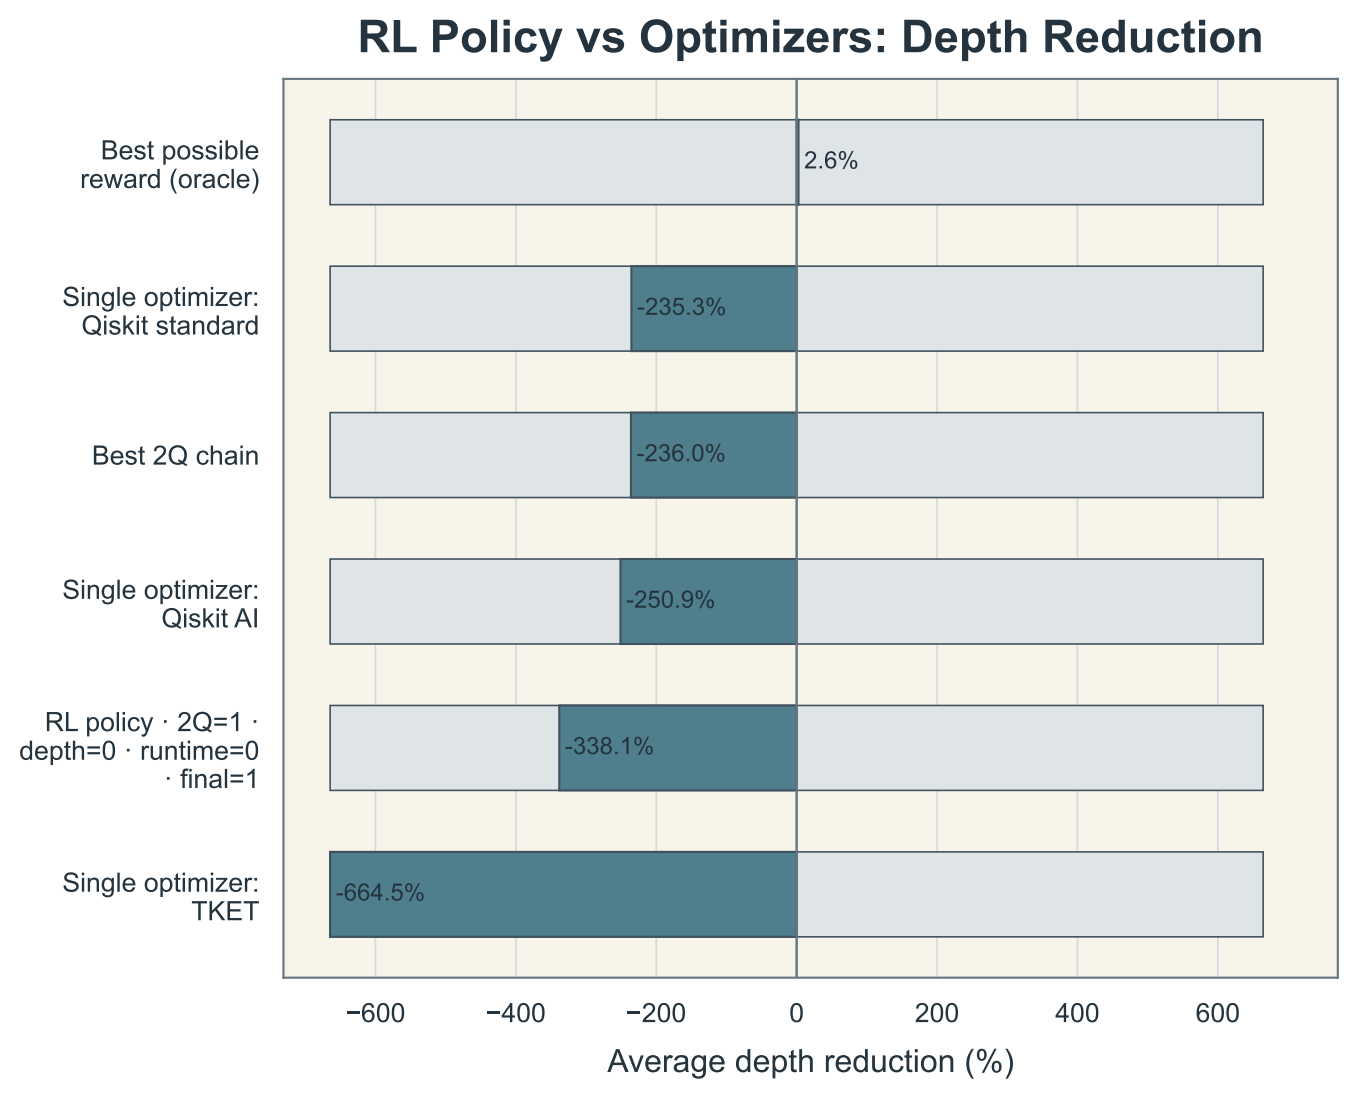
RL Policy vs Optimizers: Depth Reduction
Best possible reward (oracle)
2.6%
Single optimizer: Qiskit standard
-235.3%
Best 2Q chain
-236.0%
Single optimizer: Qiskit AI
-250.9%
RL policy · 2Q=1 · depth=0 · runtime=0 · final=1
-338.1%
Single optimizer: TKET
-664.5%
−600
−400
−200
0
200
400
600
Average depth reduction (%)